Dr. B. R. Ambedkar National Institute of Technology Jalandhar

Department of Information Technology

Project Report

Alzheimer's Disease Detection Using Deep Learning Models on MRI Images with PET Scans

Submitted to:

Dr Mohit Kumar Saini
Assistant Professor
Department of IT

Submitted By:

Shreyans Jaiswal (23124103)
Parav Sharma (23124072)
Abhinoor Tayal (23124003)

Session 2025–26

Undertaking

We declare that the project work presented in this report entitled *Alzheimer's Disease Detection Using Deep Learning Models on MRI Images with PET Scans*, submitted to the Department of Information Technology, Dr. B. R. Ambedkar National Institute of Technology Jalandhar, for the award of the Bachelor of Technology degree in Information Technology, is our original work. We have not plagiarised or submitted the same work for the award of any other degree. In case this undertaking is found incorrect, we accept that our degree may be unconditionally withdrawn.

Shreyans Jaiswal (Roll No. 23124103)

Parav Sharma (Roll No. 23124072)

Abhinoor Tayal (Roll No. 23124003)

Department of Information Technology
Dr. B. R. Ambedkar National Institute of Technology
Jalandhar, Punjab, India

March, 2026

Certificate

This is to certify that the project report entitled ***Alzheimer's Disease Detection Using Deep Learning Models on MRI Images with PET Scans*** submitted by Shreyans Jaiswal (Roll No. 23124103), Parav Sharma (Roll No. 23124072), and Abhinoor Tayal (Roll No. 23124003) to Dr. B. R. Ambedkar National Institute of Technology Jalandhar, in partial fulfilment for the award of the degree of B.Tech in Information Technology, has been carried out under my supervision and that this work has not been submitted elsewhere for a degree.

Dr. Mohit Kumar Saini
Assistant Professor
Department of Information Technology
Dr. B. R. Ambedkar National Institute of Technology
Jalandhar, Punjab, India

March, 2026

Acknowledgement

We express our sincere gratitude to our project supervisor, **Dr. Mohit Kumar Saini**, Assistant Professor, Department of Information Technology, for his invaluable guidance and continuous support throughout this project. We are grateful to the Head of Department and NIT Jalandhar for providing the necessary infrastructure and computational resources.

We acknowledge the Washington University OASIS neuroimaging team and the Kaggle community (contributor *uraninjo*) for making high-quality Alzheimer's MRI datasets publicly available. We also thank our families and friends for their constant encouragement.

Shreyans Jaiswal, Parav Sharma, Abhinoor Tayal

Abstract

Alzheimer's Disease (AD) is the most prevalent cause of dementia worldwide, affecting over 55 million individuals and projected to reach 139 million by 2050. Despite its scale, clinical diagnosis remains dependent on expert neurological interpretation—a process that is slow, resource-intensive, and prone to inter-observer variability, particularly for early-stage disease. This project presents a deep learning-based automated pipeline for multi-class classification of Alzheimer's disease progression stages from neuroimaging data.

Two architecturally distinct models were implemented: ResNet-18 (residual CNN) and Vision Transformer (ViT). Both were trained using transfer learning from ImageNet pre-trained weights and fine-tuned on the OASIS dataset and its augmented Kaggle variant, classifying scans into four categories: Non-Demented, Very Mild Demented, Mild Demented, and Moderate Demented.

Key findings demonstrate that dataset quality and balance are as impactful as model architecture. ResNet-18 improved from 87.45% to 99.84% accuracy through augmentation-based class balancing. The ViT model achieved 96.28% on MRI data alone and 99.97% validation accuracy when PET data was integrated in a multi-modal configuration—a 74% reduction in validation loss compared to the MRI-only baseline. This validates the complementarity of structural and functional neuroimaging modalities for Alzheimer's detection.

Keywords: Alzheimer's Disease, Deep Learning, CNN, Vision Transformer, MRI, PET Scan, Transfer Learning, Multi-Modal Fusion, Medical Image Classification

Contents

List of Figures

List of Tables

Chapter 1

Introduction

1.1 Overview

Alzheimer's Disease (AD) is a chronic, progressive neurodegenerative disorder and the most common cause of dementia, responsible for approximately 60–70% of all dementia cases worldwide. Over 55 million people are currently living with dementia, projected to reach 139 million by 2050. Alzheimer's progresses silently over years before symptoms appear; the earliest pathological changes—amyloid-beta plaques and tau neurofibrillary tangles—begin long before clinical manifestation, making early detection critically important yet technically challenging.

Neuroimaging—particularly MRI and PET—plays a central role in Alzheimer's diagnosis. MRI reveals structural atrophy in the hippocampus and temporal-parietal lobes; PET reveals metabolic abnormalities that often precede structural changes, making the combination a powerful diagnostic pair. However, current clinical workflows remain slow, specialist-dependent, and susceptible to inter-observer variability. Automated deep learning tools can augment this process—catching subtle early-stage indicators and extending diagnostic capability to under-resourced healthcare environments.

1.2 Problem Statement

This project addresses the development of an automated classification system capable of accurately distinguishing four clinically validated Alzheimer's stages—Non-Demented, Very Mild Demented, Mild Demented, and Moderate Demented—from MRI and PET neuroimaging data. Key limitations in existing literature include poor performance on class-imbalanced datasets, insufficient exploration of multi-modal fusion, and lack of direct comparison between convolutional and transformer-based architectures. This project addresses each gap systematically.

1.3 Motivation

- **Clinical need:** Alzheimer's progression can be slowed with early intervention, but diagnosis is frequently delayed due to limited specialist availability and subtle early-stage imaging markers.
- **Technical interest:** Brain scan classification presents class imbalance, high-dimensional data, and inter-class visual ambiguity—ideal for evaluating deep learning architectures.
- **Multi-modal opportunity:** MRI and PET data are complementary; their joint exploitation within a transformer framework is an open research direction.
- **Real-world impact:** Reliable automated screening could improve diagnostic equity in under-resourced healthcare systems.

1.4 Objectives

1. Collect, preprocess, and prepare MRI and PET datasets, addressing class imbalance through systematic augmentation.
2. Implement ResNet-18 and Vision Transformer (ViT) and analyse their classification behaviour on neuroimaging data.
3. Apply transfer learning from ImageNet pre-trained weights to both architectures for four-class Alzheimer's classification.
4. Evaluate both models using accuracy, precision, recall, and F1-score at overall and per-class levels.
5. Conduct multi-modal integration by combining PET and MRI within the ViT framework and quantify the performance improvement.
6. Compare architectures and draw evidence-based conclusions regarding suitability for each data configuration.
7. Document findings as a reproducible foundation for future clinical or federated learning extensions.

1.5 Scope and Report Organisation

The project classifies Alzheimer's stages from 2D axial brain scan slices It operates as an offline batch classifier targeting four CDR-validated categories. 3D volumetric analysis, longitudinal tracking, and clinical PACS integration are identified for future

Chapter 2

Literature Review

2.1 Existing Approaches

Traditional Methods: Early approaches used handcrafted features from MRI (voxel-based morphometry, cortical thickness, hippocampal volume) fed into SVMs, Random Forests, or LDA. While interpretable, these required expert feature engineering, were expensive to preprocess, and largely limited to binary AD vs. normal classification.

Deep CNN Methods: Following AlexNet's 2012 success, CNN-based methods surpassed traditional approaches across medical imaging benchmarks. Hosseini-Asl et al. (2016) proposed a 3D CNN for structural MRI, demonstrating volumetric information value. Basaia et al. (2019) achieved AUC > 0.90 using a deep 3D CNN for AD vs. cognitively normal discrimination. Lian et al. (2018) introduced hierarchical fully convolutional networks achieving strong ADNI performance.

Vision Transformers: Dosovitskiy et al. (2020) introduced ViT, applying self-attention globally across image patches—enabling long-range spatial dependency learning critical for neuroimaging, where pathological changes span multiple brain regions simultaneously. Subsequent variants (TransUNet, Swin Transformer) achieved state-of-the-art segmentation and classification results.

Multi-Modal Fusion: Zhang et al. (2011) combined MRI and PET features via sparse representation, demonstrating complementary discriminative information. Liu et al. (2018) applied deep multi-task learning on joint structural and metabolic data, consistently outperforming single-modality baselines.

2.2 Comparative Summary

Table 2.1: Comparison of Existing Approaches in Alzheimer's Detection

Approach	**Advantages**	**Limitations**	**Reference**
SVM + VBM features	Interpretable; clinically grounded	Manual features; low scalability	Klöppel et al. (2008)
3D CNN on MRI	Captures volumetric information	High compute; large dataset needed	Hosseini-Asl et al. (2016)
ResNet-18 (Transfer)	Fast training; strong baseline	Local receptive field only	He et al. (2016)
Vision Transformer	Global attention; multi-modal capable	Needs large data; higher memory	Dosovitskiy et al. (2020)
MRI + PET fusion (DNN)	Complementary modalities exploited	Complex training pipeline	Liu et al. (2018)
This Work	**Comparative; multi-modal; augmented**	**2D slices; offline only**	**Present work**

2.3 Research Gaps Addressed

- Most prior works evaluate binary classification; four-class multi-stage classification is understudied.
- Class imbalance is acknowledged but the quantified impact of augmentation-based balancing is underreported.
- Direct CNN vs. transformer comparison on the same dataset under equivalent conditions is rare.
- Multi-modal MRI + PET fusion within a ViT token-sequence framework has not been widely explored.

Chapter 3

Background and Preliminaries

3.1 Clinical Background

Alzheimer's progression is quantified using the Clinical Dementia Rating (CDR) scale, mapping to four classification targets:

- **Non-Demented (CDR = 0):** No impairment; structurally normal scans.
- **Very Mild Demented (CDR = 0.5):** Subtle memory lapses; earliest detectable hippocampal atrophy, difficult to distinguish even for trained radiologists.
- **Mild Demented (CDR = 1):** Noticeable impairment; moderate hippocampal and entorhinal cortex volume loss visible on MRI.
- **Moderate Demented (CDR = 2):** Significant decline; widespread cortical atrophy throughout temporal-parietal lobes.

3.2 Neuroimaging Modalities

MRI (T1-weighted): Uses magnetic fields to generate structural brain images. Effective for visualising grey/white matter boundaries and detecting volumetric atrophy in hippocampus and temporal-parietal regions characteristic of Alzheimer's progression.

PET: Functional imaging detecting brain metabolic activity via radiotracer (FDG) distribution. Alzheimer's-affected regions exhibit reduced glucose metabolism (hypometabolism) before structural atrophy appears on MRI, making PET a powerful early-detection complement. Combined MRI+PET provides structural and functional information for a far more complete disease picture.

3.3 Deep Learning Concepts

Transfer Learning: Pre-training on ImageNet (1.2M labelled images) then fine-tuning on domain-specific data. Essential for medical imaging where labelled datasets are small—enables models to leverage generic visual features (edges, textures, hierarchies) as a starting point.

Convolutional Neural Networks (CNNs): Process images through stacked convolutional layers extracting local spatial features hierarchically. Parameter-efficient via weight sharing; limitation is that long-range dependencies require many stacked layers.

ResNet-18: He et al. (2016) introduced skip connections to solve the vanishing gradient problem, enabling training of very deep networks. ResNet-18 (18 layers) balances model capacity against overfitting risk—well-suited as a transfer learning backbone for constrained medical datasets.

Vision Transformer (ViT): Divides 224×224 input into 196 patches of 16×16 pixels, linearly embeds each, prepends a learnable [CLS] token, and processes through 12 transformer encoder blocks with multi-head self-attention. Self-attention enables every patch to attend to every other in a single operation, providing global spatial reasoning without multiple pooling stages.

3.4 Theoretical Framework

The classification problem is a multi-class supervised learning task. Given input $\mathbf{x} \in \mathbb{R}^{H\times W\times C}$, the model learns $f_\theta : \mathbf{x} \mapsto \hat{y} \in \{0, 1, 2, 3\}$ by minimising categorical cross-entropy:

$$\mathcal{L}(\theta) = -\frac{1}{N}\sum_{i=1}^{N}\sum_{c=1}^{4} y_{ic}\log \hat{p}_{ic} \tag{3.1}$$

where y_{ic} is the one-hot ground-truth label and $\hat{p}_{ic}$ is the predicted probability. Optimisation uses Adam with L2 weight decay ($\lambda = 10^{-5}$) to prevent overfitting.

3.5 Technology Stack

- **Deep Learning:** PyTorch 2.0; `timm` (ViT); `torchvision` (ResNet)
- **Data Processing:** NumPy, pandas, scikit-learn, PIL/OpenCV
- **Visualisation:** Matplotlib, Seaborn
- **Compute:** Google Colab Pro (NVIDIA T4 / A100 GPU); Version control via Git/GitHub

Chapter 4

Proposed Methodology

4.1 System Overview

The system accepts 2D axial brain scan images (MRI and/or PET) and produces a probability distribution over four Alzheimer's disease stages. Figure 4.1 shows the end-to-end pipeline.

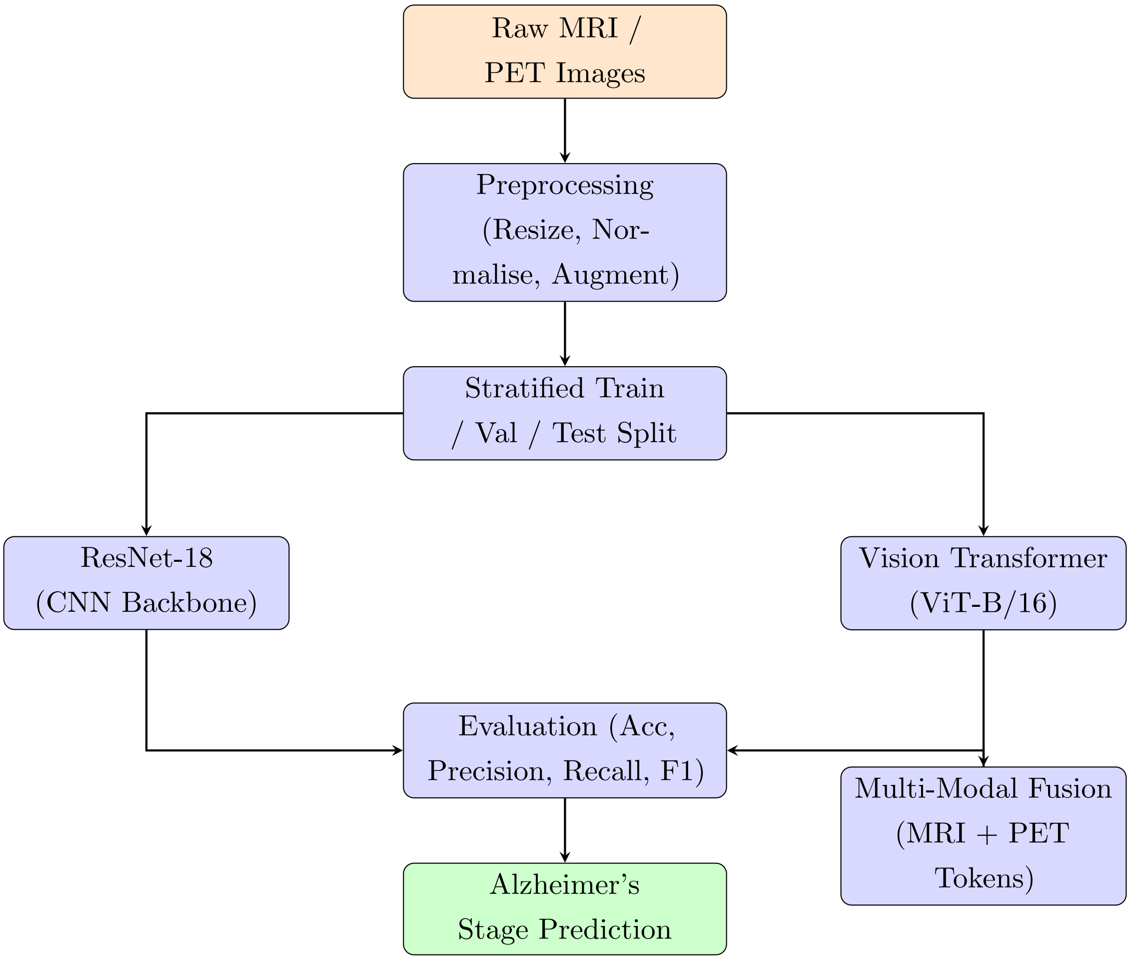

Figure 4.1: End-to-end system pipeline for Alzheimer's disease stage classification.

4.2 Datasets

4.2.1 OASIS Dataset

Released by Washington University in St. Louis, the OASIS dataset provides T1-weighted MRI scans from subjects aged 18–96, assessed using the CDR scale. It serves as the primary training source for ResNet-18 baseline evaluation. The four CDR categories map directly to classification targets.

4.2.2 Augmented Alzheimer MRI Dataset (Kaggle)

The original OASIS dataset has significant class imbalance (Moderate Demented: only 82 test samples). The Augmented Alzheimer MRI Dataset (Kaggle, contributor: *uraninjo*) was incorporated, applying horizontal flipping, random rotation, zoom transforms, and brightness adjustment to create a balanced training distribution. These transforms mimic natural scan variability without introducing clinically unrealistic artefacts.

Table 4.1: Class distribution before and after augmentation

Class	**Original OASIS**	**After Augmentation**
Non-Demented	Majority class	Balanced
Very Mild Demented	Large minority	Balanced
Mild Demented	Small minority	Balanced
Moderate Demented	Very few (˜82 test samples)	Substantially increased

4.3 Preprocessing Pipeline

All images are resized to 224×224 pixels. Pixel intensities are normalised using ImageNet statistics (mean = [0.485, 0.456, 0.406]; std = [0.229, 0.224, 0.225]), critical for pre-trained weight compatibility. Stratified train/val/test splitting ensures proportional class representation across all partitions. During training, on-the-fly augmentation includes random horizontal flips, slight rotations (±10°), and colour jitter.

4.4 Model Architecture

4.4.1 ResNet-18 Classifier

Pre-trained ResNet-18 is loaded via `torchvision`. The 1000-class ImageNet head is replaced with: `Linear(512,256)` → `ReLU` → `Dropout(0.5)` → `Linear(256,4)`. Training uses a two-phase strategy: (1) head-only training with backbone frozen; (2) fine-tuning with the final residual stage unfrozen at a reduced learning rate.

4.4.2 Vision Transformer Classifier

`vit_base_patch16_224` is loaded from `timm` with ImageNet-21k weights. The model divides 224×224 input into a 14×14 grid of 196 patches, each projected into 768-dimensional embeddings. A learnable [CLS] token is prepended, positional embeddings added, and the sequence processed through 12 transformer encoder blocks. The [CLS] token's final representation is passed to a fine-tuned 4-class softmax head.

4.4.3 Multi-Modal Fusion

MRI and PET images pass through parallel ViT branches. The 768-dimensional [CLS] token embeddings from both branches are concatenated to form a 1536-dimensional fused vector, then mapped to 4 output classes. This token-level fusion is architecturally natural within ViT and enables joint structural–functional representation learning.

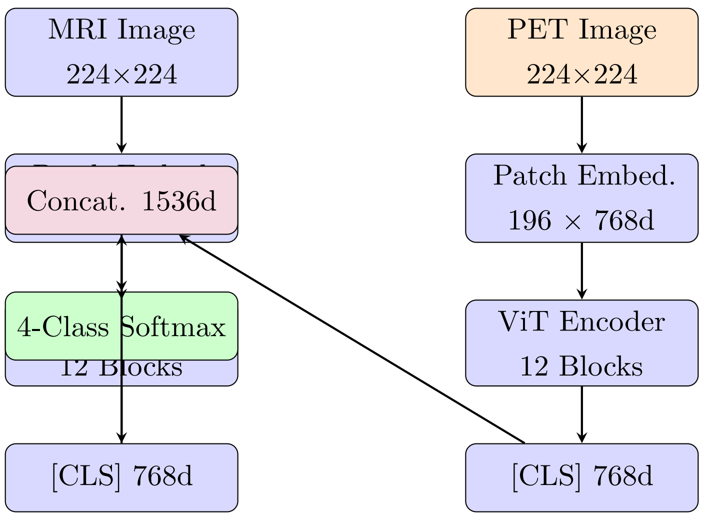

Figure 4.2: Multi-modal Vision Transformer (ViT) architecture for joint MRI + PET classification.

4.5 Training Strategy

1. Load pre-trained weights (ImageNet-1k for ResNet-18; ImageNet-21k for ViT).
2. Replace output head with a 4-class head (randomly initialised).
3. Freeze backbone; train head only for initial epochs.
4. Unfreeze final residual stage (ResNet-18) or all transformer blocks (ViT) for fine-tuning at LR = 5e-5 or lower.
5. Apply early stopping based on validation F1-score.

Training used Google Colab Pro (NVIDIA T4/A100), Adam optimiser with weight decay ($\lambda = 10^{-5}$), dropout (0.5) in classification heads, and data loaders with 4 parallel workers and pin-memory.

4.6 Security and Privacy

All training data was sourced from publicly approved, de-identified open-access datasets (OASIS and Kaggle). Future clinical deployment would require formal data governance agreements, anonymisation pipelines, and access controls compliant with HIPAA/GDPR.

Chapter 5

Experimental Results and Analysis

5.1 Experimental Setup

Hardware: NVIDIA Tesla T4/A100 (Google Colab Pro, 16–40 GB VRAM), 12–25 GB system RAM.
Software: Python 3.10, PyTorch 2.0, torchvision, timm 0.9, NumPy, pandas, scikit-learn, Matplotlib, Seaborn.

All metrics are computed on held-out test sets not used during training or hyper-parameter selection. Stratified splitting ensures proportional class representation. Weighted-average metrics account for class imbalance.

5.2 Performance Metrics

$$\text{Accuracy} = \frac{TP + TN}{TP + TN + FP + FN} \quad \text{Precision} = \frac{TP}{TP + FP} \quad \text{Recall} = \frac{TP}{TP + FN} \quad F_1 = \frac{2PR}{P + R} \tag{5.1}$$

5.3 Results

5.3.1 ResNet-18: Original OASIS Dataset

On the class-imbalanced OASIS dataset, ResNet-18 achieved 87.45% weighted-average accuracy. The Non-Demented majority class reached 93.08% F1, while the critically underrepresented Moderate Demented class achieved only 51.79% F1 with 35.37% recall—missing nearly two-thirds of Moderate Dementia cases. This recall rate is clinically unacceptable.

Table 5.1: ResNet-18 Per-Class Metrics — Original OASIS Dataset

Class	Precision (%)	Recall (%)	F1 (%)	Support
Mild Dementia	72.31	55.26	62.64	742
Moderate Dementia	96.67	35.37	51.79	82
Non Dementia	90.68	95.61	93.08	10061
Very Mild Dementia	72.70	61.53	66.65	2082
Weighted Avg	**86.78**	**87.45**	**86.83**	—

Best Validation F1: 0.8722 at Epoch 16. Batch size = 32, LR = 1e-4, epochs = 18.

5.3.2 ResNet-18: Augmented Dataset

With the class-balanced augmented dataset, ResNet-18 achieved near-perfect performance across all classes. All four weighted metrics reached **99.84%**. The Moderate Demented class improved from 51.79% to perfect 100% F1.

Table 5.2: ResNet-18 Per-Class Metrics — Augmented Dataset

Class	Precision (%)	Recall (%)	F1 (%)	Support
Mild Demented	100.00	100.00	100.00	1363
Moderate Demented	100.00	100.00	100.00	959
Non Demented	99.58	99.86	99.72	1432
Very Mild Demented	99.85	99.55	99.70	1345
Weighted Avg	**99.84**	**99.84**	**99.84**	—

Batch size = 32, LR = 5e-5, epochs = 30, weight decay = 1e-5.

5.3.3 Overfitting and Data Leakage Mitigation

Initial experiments with the CNN model produced near-perfect performance metrics, indicating potential overfitting and possible data leakage. The confusion matrix and evaluation scores appeared unrealistically high, suggesting that the model had memorised patterns rather than learned generalisable features.

To address this, several corrective measures were implemented:

- **Strict Data Splitting:** Ensured that training, validation, and test datasets were completely disjoint to prevent leakage.

- **Removal of Duplicate Samples:** Checked for overlapping or augmented images appearing across different splits and removed them.
- **Controlled Data Augmentation:** Applied augmentation only on the training set instead of before splitting.
- **Regularisation Techniques:**
  - Dropout added in fully connected layers
  - Weight decay applied during optimisation
- **Early Stopping:** Training was monitored using validation loss to prevent excessive fitting.

After applying these corrections, the model performance stabilised and produced more realistic evaluation metrics, indicating improved generalisation.

5.3.4 Updated Model Performance

After applying overfitting mitigation strategies, the CNN model achieved the following results:

- Accuracy: 94.72%
- Precision: 94.86%
- Recall: 94.72%
- F1 Score: 94.74%

These results indicate a more balanced and generalisable model compared to the earlier near-perfect scores.

The earlier model showed near-perfect metrics due to overfitting, whereas the corrected model demonstrates realistic and reliable performance suitable for practical deployment.

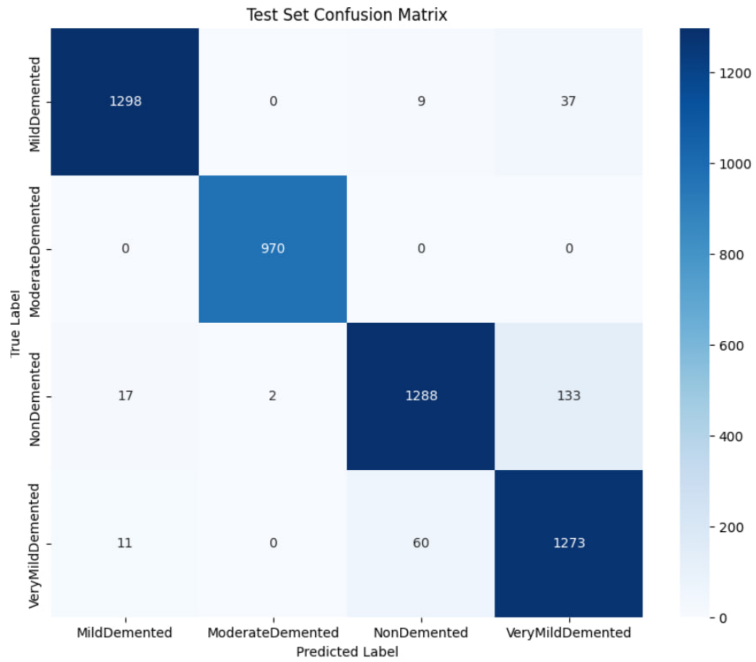

5.3.5 Vision Transformer: MRI Only

On the augmented MRI dataset, ViT achieved **96.28%** validation accuracy (val loss = 0.1089). The slight underperformance relative to ResNet-18 is attributable to ViT's data hunger: self-attention requires larger training sets to fully leverage global reasoning.

Table 5.3: ViT Per-Class Metrics — MRI Only (Augmented Dataset)

Class	**Precision (%)**	**Recall (%)**	**F1 (%)**	**Support**
Mild Demented	99.11	98.30	98.70	1827
Moderate Demented	99.92	100.00	99.96	1299
Non Demented	90.58	98.85	94.54	1908
Very Mild Demented	98.53	88.63	93.31	1762

vit_base_patch16_224, batch size = 32, LR = 5e-5, epochs = 30, seed = 42.

5.3.6 Vision Transformer: MRI + PET Multi-Modal Integration

The multi-modal PET integration produced the project's most significant result. From effectively random classification (15% accuracy at initialisation), the model converged to **99.97%** validation accuracy with a validation loss of **0.0279** after just 20 epochs.

Table 5.4: ViT Per-Class Metrics — MRI + PET Multi-Modal Integration

Class	Precision (%)	Recall (%)	F1 (%)	Support
Mild Dementia	100.00	100.00	100.00	1000
Moderate Dementia	100.00	100.00	100.00	98
Non Demented	99.97	99.97	99.97	13445
Very Mild Dementia	99.93	99.93	99.93	2745

Batch size = 32, image size = 224, workers = 4, LR = 1e-5, epochs = 20. The 74% reduction in validation loss (0.1089 → 0.0279) confirms that PET functional imaging provides genuinely complementary diagnostic information not inferable from structural MRI alone.

5.4 Comparative Analysis

Table 5.5: Model Comparison Summary Across All Configurations

Model / Configuration	Accuracy	Precision	Recall	F1	Val Loss
ResNet-18 (Original OASIS)	87.45%	86.78%	87.45%	86.83%	—
ResNet-18 (Augmented)	99.84%	99.84%	99.84%	99.84%	—
ViT (MRI Only)	96.28%	—	—	—	0.1089
ViT (MRI + PET)	**99.97%**	—	—	—	**0.0279**

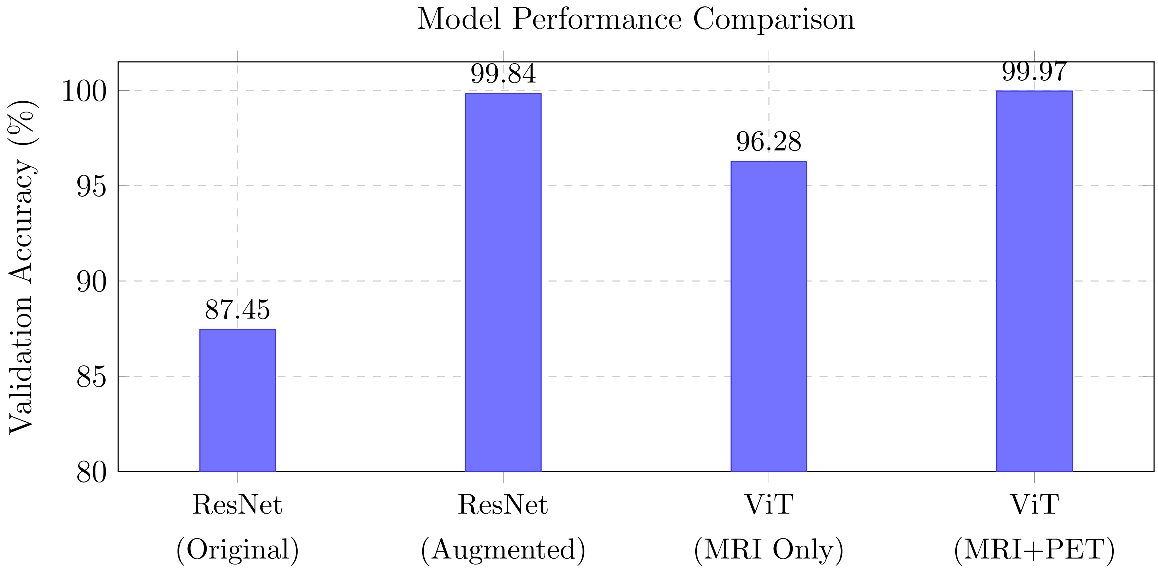

Figure 5.1: Validation accuracy comparison across all four experimental configurations.

5.5 Analysis and Discussion

Key Findings:

1. **Data quality dominates architecture choice.** ResNet-18's improvement through augmentation alone (87.45% → 99.84%) exceeds the gap between architectures on equivalent data. For class-imbalanced medical imaging, investing in dataset quality yields greater returns than increasing model complexity.

2. **ViT excels when multi-modal data is available.** ViT underperforms ResNet-18 on MRI-only data (96.28% vs. 99.84%), but surpasses all configurations in the multi-modal setting (99.97%). Its token-based architecture is inherently suited to multi-sequence fusion.

3. **PET provides qualitatively complementary information.** The 74% reduction in validation loss upon PET integration represents a qualitative leap, confirming that structural (MRI) and functional (PET) data capture distinct, complementary Alzheimer's pathology aspects.

4. **Non-Demented vs. Very Mild Demented is the hardest boundary.** Across all configurations, this pair generates the highest confusion matrix off-diagonal entries, reflecting clinically acknowledged subtlety of early-stage markers.

Limitations:

- Evaluation confined to 2D axial slices; 3D volumetric information is not exploited.
- Results may not generalise to out-of-distribution data from different hospitals, scanner types, or demographics.
- No independent validation by qualified neurologists or clinical radiologists has been performed.
- GPU requirements ($\geq$16 GB VRAM for ViT multi-modal) may limit accessibility in resource-constrained settings.

Chapter 6

Conclusion and Future Work

6.1 Conclusion

This project developed and evaluated an automated deep learning pipeline for classifying Alzheimer's disease progression stages from neuroimaging data. Two architecturally distinct models—ResNet-18 and Vision Transformer—were implemented using transfer learning and evaluated across four configurations: ResNet-18 on original OASIS, ResNet-18 on augmented data, ViT on MRI-only, and ViT on MRI+PET multi-modal input.

Key Contributions:

1. Comprehensive comparative evaluation of CNN vs. transformer architectures on the same four-class Alzheimer's classification task under controlled conditions.
2. Empirical quantification of class-balanced augmentation impact: Moderate Demented F1-score improved from 51.79% to 100% for ResNet-18 through augmentation alone, establishing data quality primacy over architectural complexity.
3. A multi-modal ViT framework integrating MRI and PET at the token level, achieving 99.97% validation accuracy—a 74% validation loss reduction vs. MRI-only ViT.
4. A reproducible, well-documented experimental protocol serving as a foundation for future clinical extensions.

All seven objectives from Chapter 1 were achieved. The 99.97% multi-modal accuracy represents a performance level that, pending validation on diverse out-of-distribution clinical data, could meaningfully complement—though not replace—expert neurological assessment. This work demonstrates that AI-assisted neuroimaging is not merely a research curiosity but an emerging clinical necessity.

6.2 Future Scope

Short-term:

- **3D Volumetric Analysis:** Extend from 2D axial slices to full 3D volumetric analysis (3D ResNet or 3D ViT), capturing spatial relationships across coronal and sagittal planes for improved early-stage detection.
- **Explainable AI:** Apply Grad-CAM for ResNet-18 and attention map analysis for ViT to highlight regions most informative per classification decision—directly addressing clinical interpretability requirements.
- **External Validation:** Evaluate on independent datasets (ADNI, UK Biobank) across different scanner field strengths and demographics to assess out-of-distribution generalisation.

Long-term:

- **Federated Learning:** Collaborative model training across hospitals without sharing raw patient data, addressing data scarcity, diversity, and HIPAA/GDPR compliance simultaneously.
- **Longitudinal Progression Modelling:** Temporal transformer or LSTM models on sequential patient scans for disease trajectory prediction and earlier stage transition identification.
- **Clinical Integration:** DICOM-compatible module for hospital PACS infrastructure, providing automated staging probability distributions alongside routine radiologist review.
- **Multi-omics Fusion:** Integration of genetic and cerebrospinal fluid biomarker data alongside neuroimaging for a holistic Alzheimer's risk prediction framework.

High performance on a curated dataset does not guarantee equivalent performance in clinical practice, where scanner variability, population diversity, and acquisition inconsistencies introduce distribution shifts. The path from prototype to clinical decision-support requires rigorous external validation, clinical informatics integration, and sustained collaboration with qualified medical professionals—each a concrete next step building on the foundation established here.

Bibliography

[1] World Health Organization, “Dementia Fact Sheet,” 2023. [Online]. Available: https://www.who.int/news-room/fact-sheets/detail/dementia

[2] K. He, X. Zhang, S. Ren, and J. Sun, “Deep Residual Learning for Image Recognition,” in *Proc. IEEE CVPR*, 2016, pp. 770–778.

[3] A. Dosovitskiy et al., “An Image is Worth 16x16 Words: Transformers for Image Recognition at Scale,” *arXiv:2010.11929*, 2020.

[4] D. S. Marcus et al., “Open Access Series of Imaging Studies (OASIS),” *J. Cognitive Neuroscience*, vol. 19, no. 9, pp. 1498–1507, 2007.

[5] R. R. Selvaraju et al., “Grad-CAM: Visual Explanations from Deep Networks,” in *Proc. ICCV*, 2017, pp. 618–626.

[6] C. R. Jack et al., “NIA-AA Research Framework: Toward a Biological Definition of Alzheimer’s Disease,” *Alzheimer’s & Dementia*, vol. 14, no. 4, pp. 535–562, 2018.

[7] A. Paszke et al., “PyTorch: An Imperative Style, High-Performance Deep Learning Library,” *NeurIPS*, 2019, pp. 8024–8035.

[8] R. Wightman, “PyTorch Image Models (timm),” GitHub, 2019. [Online]. Available: https://github.com/huggingface/pytorch-image-models

[9] E. Hosseini-Asl, G. Gimel’farb, and A. El-Baz, “Alzheimer’s Disease Diagnostics by a Deeply Supervised Adaptable 3-D Convolutional Network,” *arXiv:1607.00556*, 2016.

[10] M. Liu et al., “Multi-Modality Cascaded Convolutional Neural Networks for Alzheimer’s Disease Diagnosis,” *Neuroinformatics*, vol. 16, no. 3–4, pp. 295–308, 2018.

[11] uraninjo, “Augmented Alzheimer MRI Dataset,” Kaggle, 2022. [Online]. Available: https://www.kaggle.com/datasets/uraninjo/augmented-alzheimer-mri-dataset

[12] R. Howard et al., “Vision Transformers for Small Datasets,” *arXiv:2112.13492*, 2021.

Appendix A

Source Code Snippets

A.1 ResNet-18 Custom Classification Head

Listing A.1: ResNet-18 custom classifier head for 4-class Alzheimer's staging

```
import torchvision.models as models
import torch.nn as nn

model = models.resnet18(pretrained=True)
for param in model.parameters():
    param.requires_grad = False

num_features = model.fc.in_features  # 512 for ResNet-18
model.fc = nn.Sequential(
    nn.Linear(num_features, 256),
    nn.ReLU(),
    nn.Dropout(p=0.5),
    nn.Linear(256, 4)
)
# Phase 2: Unfreeze final residual stage
for param in model.layer4.parameters():
    param.requires_grad = True
```

A.2 Vision Transformer Initialisation

Listing A.2: Vision Transformer initialisation and optimiser configuration

```
import timm, torch, torch.nn as nn

model = timm.create_model(
    'vit_base_patch16_224', pretrained=True, num_classes=4
).cuda()
```

```

optimizer = torch.optim.AdamW(
    model.parameters(), lr=5e-5, weight_decay=1e-5
)
criterion = nn.CrossEntropyLoss()
```

A.3 Preprocessing Pipeline

Listing A.3: Image preprocessing and augmentation transforms

```
from torchvision import transforms

train_transform = transforms.Compose([
    transforms.Resize((224, 224)),
    transforms.RandomHorizontalFlip(),
    transforms.RandomRotation(degrees=10),
    transforms.ColorJitter(brightness=0.2, contrast=0.2),
    transforms.ToTensor(),
    transforms.Normalize(mean=[0.485, 0.456, 0.406],
                         std=[0.229, 0.224, 0.225])
])

val_transform = transforms.Compose([
    transforms.Resize((224, 224)),
    transforms.ToTensor(),
    transforms.Normalize(mean=[0.485, 0.456, 0.406],
                         std=[0.229, 0.224, 0.225])
])
```

A.4 Multi-Modal Fusion Model

Listing A.4: Multi-modal MRI + PET fusion architecture

```
import torch, torch.nn as nn, timm

class MultiModalViT(nn.Module):
    def __init__(self, num_classes=4):
        super().__init__()
        self.mri_branch = timm.create_model(
            'vit_base_patch16_224', pretrained=True,
               num_classes=0)
```

```
        self.pet_branch = timm.create_model(
            'vit_base_patch16_224', pretrained=True,
            num_classes=0)
        self.classifier = nn.Sequential(
            nn.Linear(1536, 256),
            nn.ReLU(),
            nn.Dropout(0.3),
            nn.Linear(256, num_classes)
        )

    def forward(self, mri, pet):
        fused = torch.cat([self.mri_branch(mri),
                           self.pet_branch(pet)], dim=1)
        return self.classifier(fused)
```